

Five Seat Protocol

How to run your five-helper answer checker, one box at a time.

1

OpenAI

2

Google

3

Mistral

4

xAI

5

Anthropic

Five helpers read the same page. Each one sits in its own room and cannot see what the others wrote. Then your computer checks their work by doing the maths itself. Wrong answers get crossed out. Whatever is still standing at the end is your answer.

FILL THIS IN BEFORE YOU START

YOUR NAME

TODAY'S DATE

RUN NUMBER

Tick the boxes as you go. If a box will not tick honestly, stop there and fix it. Skipping a box is how people spend money on a run that cannot answer anything.

THREE THINGS TO KNOW BEFORE YOU START



Most of this is free.

Sections 01, 04, 05 and 06 cost nothing at all. Only two steps in the whole booklet touch a paying service.



Say hello before you spend.

Section 02 sends five tiny messages for well under a cent. It is the cheapest way to find out something is broken.



One broken helper ruins the whole run.

If any of the five fails, the run cannot measure anything -- but the other four still charge you.

00 What this thing actually does

READ ME FIRST

Imagine you have a question, and two or three possible answers. You are not sure which answer is right.

You give the question to five different AI helpers. We call them seats. Seat 1 is OpenAI, seat 2 is Google, and so on. Each seat is in its own room. No seat can see what another seat wrote. That part matters more than anything else here: two people who peek at each other's homework stop being two people.

Each seat writes down small statements called claims. A claim is a short thing that is either true or false, like "3 plus 1 equals 4". Then your computer checks those claims by itself, by doing the maths. It does not ask an AI whether the AI was right.

Any possible answer that leaned on a claim that failed gets crossed out. What is left standing is your answer.

A REAL EXAMPLE

Your note says	We sold 3 in October and 1 in November, so the quarter total is 5.
Answer A	The total is 4.
Answer B	The total is 5, like the note says.
Computer does	3 + 1, and gets 4.
Result	Answer B is crossed out. Answer A is left standing. Nobody voted. The sum decided.



Why five helpers and not one?

One helper gives you one opinion. Five helpers who all make the same mistake give you one mistake, five times over -- and it looks like agreement, which is worse than being wrong on your own. The run measures whether your five really do think differently. That measurement is the whole point of doing this.

00b The colours and shapes on these pages

MARK	MEANING	WHAT TO DO
GREEN	Free. Costs nothing at all.	Do all of it. There is no reason to skip a free check.
YELLOW	Costs a tiny amount of money.	Fine to do. Under one cent.
RED	STOP. Something is broken.	Do not go past this until it is fixed. Going on costs money and teaches you nothing.
[]	A box you tick.	Click it in this PDF and it stays ticked when you save.
—	A line you type on.	Click and type. It saves with the file.

NOTES — anything you noticed, anything that felt wrong

What is in this booklet

00	What this thing actually does The idea, in plain words	2
01	Warm-up checks Free. Nothing can cost money here.	4
02	Say hello to all five Tiny cost. The stop gate.	5
03	The real run This one costs money.	6
04	Read what came back The number that matters is not the answer.	7
05	If a helper did not answer Look up the error. Do not guess.	8
06	Finish and sign Write down what it cost.	9

**Do these in order, top to bottom**

The order is not a suggestion. Section 02 is a gate: if all five helpers do not say hello, section 03 will take your money and give you nothing back. That has already happened twice on this project.

Where everything lives

Everything happens in one folder. To get there, open Terminal (press Command and Space together, type Terminal, press Enter) and paste this line, then press Enter:

**Go to the folder. Do this once, every time you open Terminal.**

If a command later says "no such file", you are in the wrong folder. Paste this line again.

```
cd ~/forge-talent-connections/adjudication
```

01 Warm-up checks

FREE - \$0.00

Nothing in this section touches an AI company. Nothing here can cost you a cent. Do every single one.

☐ Check the right Python is being used.

You want it to say 3.13 or higher. If it says 3.9, stop -- the tool cannot work on that and needs setting up again.

```
.venv/bin/python --version
```

WHAT IT SAID

☐ Check all five keys are in place.

A key is like a password that lets you use each AI company. You want it to say 5 of 5 keys filled and No vendor mismatches detected. This never shows your keys on screen.

```
./set-key.command check
```

HOW MANY FILLED

ANY MISMATCH?

☐ Check the settings file is healthy.

You want the words PROFILES OK. If it lists problems instead, it tells you every single one at once, so you only have to go round the loop once.

```
.venv/bin/python run_adjudication.py --check-profiles profiles.json
```

OK OR PROBLEMS?

☐ Run the pretend version.

This uses fake helpers, so it is free. You want to see SURVIVOR: c_true and a section headed HOLES. That proves the machine works before you spend anything.

```
.venv/bin/python run_adjudication.py --demo
```

SAW SURVIVOR?

SAW HOLES?

☐ Run the self-tests.

You want 503 passed. If any test fails, something is broken in the tool itself, not in your setup.

```
.venv/bin/python -m pytest test_suite.py test_properties.py -q
```

NUMBER PASSED

NOTES — anything you noticed, anything that felt wrong

**Check your secrets cannot escape.**

The first command should NOT list .env or profiles.json. The second should name a rule for each. Those two files hold your keys.

```
git status --short
git check-ignore -v .env profiles.json
```

NOTES — anything you noticed, anything that felt wrong

02 Say hello to all five

TINY COST

This is the smallest paid step that exists. It sends each helper six words and asks it to reply with one word: OK. Five tiny messages, all five together costing less than one cent.

**Say hello to every helper.**

You can also do just one, like: `diagnose-seats.py seat_1`

```
.venv/bin/python diagnose-seats.py \
  seat_1 seat_2 seat_3 seat_4 seat_5
```

**Write down what each one said.**

You are looking for two things on every seat: HTTP 200, and `text_path` resolved to: 'OK'.

1 OPENAI

2 GOOGLE

3 MISTRAL

4 XAI

5 ANTHROPIC

**STOP GATE -- do not walk past this one**

All five must say 200 AND show the word OK. If even one does not, go to section 05 and fix it first. Here is why it matters so much: if any single helper fails during the real run, the whole run becomes unmeasurable -- but the other four still charge you. You pay full price for an answer that cannot tell you anything.

**A 200 on its own is not good enough**

A helper can reply 200, meaning "all fine", while the actual words sit somewhere the tool is not looking. The tool then records that helper as having read your page and found nothing to say. That is the most expensive kind of mistake here, because it does not look like a mistake. It looks like data.

NOTES — anything you noticed, anything that felt wrong

03 The real run

COSTS MONEY

About 25 messages go out. Replace ARTIFACT.txt with your page, and give the audit file a number you have never used before.

☐ Start the run.

The audit file is your permanent receipt of what happened. Never reuse a name -- writing over it destroys the record.

```
.venv/bin/python run_adjudication.py ARTIFACT.txt \  
--profiles profiles.json \  
--candidates candidates.json \  
--audit run-003.jsonl
```

ARTIFACT USED

AUDIT FILE NAME

☐ Check the very first line.

It should say: env: loaded /Users/.../adjudication/.env . If it says no .env found, stop -- it cannot see your keys.

☐ Stay and watch until the first pass finishes.

Helper 2 (Google) has no way to limit how much it writes, so if it runs long it keeps charging while you are not looking. Once pass 1 finishes with no errors, it is safe to walk away.

☐ Do not shut the laptop and do not lose Wi-Fi.

If the internet drops, every helper fails at once and the rest of the run is wasted. You still pay for the messages already sent.



It will look frozen. It is not.

The screen shows one line and then nothing for several minutes. That is normal: it prints nothing until a whole pass has finished, and a pass means all five helpers answering. Leave it alone.

NOTES — anything you noticed, anything that felt wrong

04

Read what came back

FREE - \$0.00

The survivor line is the least interesting part of the report. Read these three things, in this order.



First: did every helper answer?

Look for the words SEAT ERROR. If you see them, the report now tells you the reason right underneath. Any seat error means that pass does not count towards the measurement.

ANY SEAT ERRORS?



Second: find the number called rho.

It sits under the heading CONVERGENCE: ELIMINATIVE, OR COLLAPSE? If it says NOT MEASURABLE, the run told you nothing about whether your helpers think differently -- treat the answer as unproven.

RHO

COVERAGE



Third: read the HOLES section.

Holes are the parts the run could not settle. This is part of the answer, not an apology. A clean run still has holes.

SURVIVOR

HOW MANY HOLES



If all five agreed exactly, be suspicious

On a real question, five helpers thinking for themselves do not produce identical lists. If they did, they probably share the same blind spot -- so your panel of five is really worth about one. Agreement is not proof. Sometimes it is the opposite.



About sources and web links

Nothing in this tool checks whether a source actually exists. It can tell that something is shaped like a proper source, but not that it is real. So any source a helper gives you must be opened and checked by you, with your own eyes, before you use it.

NOTES — anything you noticed, anything that felt wrong

05 If a helper did not answer

FREE TO LOOK UP

Find the error you saw in the left column. Every one of these has happened for real on this machine and been fixed.

WHAT YOU SAW	WHAT IT REALLY MEANS	WHAT TO DO
401	The key is wrong, old, or was cancelled. This is NOT about money.	Make a brand new key on that company's website, then run <code>./set-key.command</code> and the seat number.
429 quota	The key is fine. There is no money on the account, or the free plan gives you zero of this model.	Turn on billing for that company. Careful: a ChatGPT Plus subscription is a different thing and gives you none of this.
400	You asked for a setting that helper does not allow.	Look in <code>profiles.json</code> for that seat and take the setting out.
no text at the configured path	It answered fine, but the tool looked in the wrong place for the words.	The reply path for that seat needs fixing in <code>profiles.json</code> .
URL error / nodename	Your internet dropped. Nothing is broken.	Reconnect and run it again.
panel incomplete	A key exists for a seat with no settings block.	Add that seat to <code>profiles.json</code> .

**Never guess which one it is**

On this project three different AI systems each gave a confident, well-argued explanation for why these helpers were failing. All three were wrong, every time. The only thing that settled it was running the hello test in section 02 and reading what the company actually said. It costs a fifth of a cent. Guessing costs a whole run.

NOTES — anything you noticed, anything that felt wrong

06 **Finish and sign**

FREE - \$0.00

☐ **Write down what it cost and what it produced.**

MESSAGES SENT

MONEY SPENT

AUDIT FILE

☐ **Deal with the leftover questions.**

The run puts unanswered things in a queue. Answer them yourself, then run it again with --adjudications so your decisions get folded in. An answer with an open queue is a favourite, not a conclusion.

```
.venv/bin/python run_adjudication.py ARTIFACT.txt \
  --profiles profiles.json --candidates candidates.json \
  --adjudications decisions.json --audit run-003b.jsonl
```

☐ **Check nothing secret leaked.**

Your keys live in .env and your settings in profiles.json. Neither should ever appear in this list.

```
git status --short
```

SIGN OFF

DID ONE ANSWER SURVIVE?

WAS RHO MEASURABLE?

HOW MANY HOLES ARE LEFT?

YOUR SIGNATURE

☒ **You are done**

A run that ends with one survivor, a real rho, and a short list of honest holes is a good run. A run that ends with NOT RESOLVED and holes you can act on is also a good run. The only bad run is one that reports certainty it did not earn.

NOTES — anything you noticed, anything that felt wrong